

TinySepsis: Calibrated, Low-Resource Early Warning for Sepsis with Compact Temporal Models

Elían Alfonso López Preciado*
Independent Researcher, Mexico

August 2026

Abstract

Sepsis remains one of the leading causes of in-hospital mortality worldwide, and the marginal benefit of early antibiotic administration and fluid resuscitation makes early warning systems clinically valuable. However, most published sepsis-prediction models are evaluated primarily on discrimination (AUROC), trained and deployed on infrastructure unavailable to low-resource hospitals, and rarely audited for probability calibration or false-alarm burden — two properties that determine whether a warning system is trusted or ignored in practice. We present **TinySepsis**, a compact (under 200K trainable parameters) gated recurrent unit (GRU) model for 4-, 6-, and 8-hour early prediction of sepsis onset from routinely collected vital signs and laboratory values, explicitly engineered for training and inference on a single consumer GPU (8 GB VRAM) or CPU-only hardware. We combine (i) explicit missingness encoding (observation masks and time-since-last-measurement, rather than naive imputation), (ii) post-hoc probability calibration (temperature scaling and isotonic regression), and (iii) conformal risk control to target an empirical false-alarm-rate budget on held-out data — a marginal, exchangeability-dependent guarantee whose erosion under real cross-hospital shift we measure and report rather than assume away. Using the open-access PhysioNet/Computing in Cardiology Challenge 2019 dataset [Reyna et al., 2019b], we train on one hospital system and evaluate, without any further tuning, on a second hospital system from the same data-collection effort but never seen during development — a genuine, uncurated cross-institution stress test, though not a substitute for a fully independent-cohort external validation against e.g. MIMIC-IV or eICU, which we did not have credentialed access to (Limitations). On the internal test set, gradient-boosted tree baselines (XGBoost, LightGBM) and logistic regression outperform TinySepsis on raw discrimination (AUROC 0.77–0.81 vs. 0.74, mean over 5 seeds) and continue to do so on the second hospital as well; XGBoost remains the best performer there too (0.686). What changes sharply is the *size of the drop*: every tabular baseline loses 0.115–0.126 AUROC on average moving to the second hospital, while TinySepsis loses only 0.072 – roughly 40% smaller than any tabular baseline’s average loss – and, unlike the single-run numbers a first pass at this analysis would report, this gap survives replication across 5 random seeds per model (TinySepsis’s degradation is smaller than *every one* of the 15 baseline-seed runs; Welch’s t $p \leq 0.002$ and Mann-Whitney $p = 0.004$ against each baseline, Bonferroni-corrected, Cohen’s $d \in [-4.9, -3.6]$). A second, unanticipated result complicates this picture: at each model’s own conformal-calibrated alarm threshold, TinySepsis is the *only* learned model whose clinical utility turns negative on the second hospital, because its calibrated false-alarm rate exceeds target there while the tabular baselines’ does not – so TinySepsis’s AUROC ranking transfers across institutions better than the baselines’, but its calibrated decision threshold transfers worse, even though the margin by which it does so is modest (external utility -0.023 , essentially break-even, vs. $+0.12$ to $+0.16$ for the tabular baselines). Both findings are limited to a single hospital pair from one

*Correspondence: inovation.of.the.pear@gmail.com. Code, trained weights, and reproduction scripts: <https://github.com/elianalfonsolopezpreciado/tinysepsis> (public, MIT-licensed).

data-collection effort (Limitations). We report all of this alongside calibration, decision-curve net benefit, alarms-per-1000-patient-hours, and mean lead time before clinical suspicion, and an efficiency profile (parameter count, ONNX export size, CPU/GPU latency) intended to make deployability auditable rather than asserted. All code, trained weights, and the exact data pipeline are released for independent reproduction. TinySepsis is a research prototype, not a validated medical device, and is not intended for clinical deployment without prospective validation and regulatory review.

Keywords: sepsis, early warning systems, clinical machine learning, calibration, conformal prediction, low-resource deployment, missing data, external validation

Contents

1	Introduction	4
1.1	Contributions	4
1.2	What this paper is not	5
2	Related Work	5
3	Problem Formulation	6
4	Dataset	6
5	Model Architecture	7
6	Calibration and Conformal Risk Control	8
7	Clinical Utility Metric	8
8	Baselines	9
9	Experiments and Results	9
9.1	Internal test set (Hospital A, held-out patients)	9
9.2	Cross-institution evaluation (Hospital B)	10
9.3	Multi-seed replication	10
9.4	Calibration and conformal risk control	11
9.5	Efficiency and deployability	13
9.6	Ablations	13
9.7	Feature attribution	14
9.8	Subgroup analysis	14
10	Discussion	15
11	Limitations	16
12	Ethical Considerations	17
13	Local Deployment and Reproducibility	17
14	Conclusion	17

1 Introduction

Sepsis — a dysregulated host response to infection that leads to organ dysfunction [Singer et al., 2016] — affects an estimated 1.7 million adults annually in the United States alone [Rhee et al., 2017] and is present in nearly one in three in-hospital deaths [Rhee et al., 2019]. Each hour of delay in antibiotic administration after sepsis is recognized is associated with increased mortality [Seymour et al., 2017], which has motivated a large body of work on automated, EHR-integrated early warning systems intended to flag at-risk patients before a clinician would otherwise recognize them.

Despite this activity, the translational track record of sepsis prediction models is mixed. The most consequential recent case study is the Epic Sepsis Model, a proprietary model embedded in electronic health record software at hundreds of U.S. hospitals: an external validation at a large academic medical center found an AUROC of only 0.63, well below the vendor’s reported performance, along with a sensitivity of 33% at the deployed operating point [Wong et al., 2021, Habib et al., 2021]. This gap between internal and external performance is not an isolated failure; it reflects a structural pattern in clinical machine learning where models tuned and validated on a single institution’s case mix, documentation habits, and monitoring equipment degrade when deployed elsewhere [Rajkomar et al., 2019, Obermeyer et al., 2019]. A second, related failure mode is alarm fatigue: warning systems that fire frequently and imprecisely are downweighted or ignored by clinical staff regardless of their nominal sensitivity, so a system’s operating characteristics under a fixed false-alarm budget matter as much as its area under the ROC curve [Sendak et al., 2020].

We argue that three properties, largely orthogonal to raw discriminative performance, determine whether a sepsis early-warning system is trustworthy and deployable, particularly in resource-constrained settings that cannot rely on continuous cloud connectivity or GPU clusters:

1. **Calibration.** A predicted risk of 80% should correspond, empirically, to an approximately 80% observed event rate in the relevant population. Uncalibrated risk scores are common in deep learning [Guo et al., 2017] and are actively dangerous in a clinical triage context, where the number itself (not just its rank) is presented to a decision-maker [Van Calster et al., 2019].
2. **Bounded false-alarm rate.** Rather than reporting sensitivity at an arbitrarily chosen threshold, a deployable system should offer an explicit, statistically grounded control on how often it is wrong when it alarms, ideally with a finite-sample guarantee that survives distribution shift as well as possible [Angelopoulos et al., 2022, Vovk et al., 2005].
3. **Deployability.** The system should run on hardware realistically available in a lower-resource hospital or a rural ICU: no proprietary cloud API, no multi-GPU cluster, and ideally CPU inference latency compatible with hourly EHR polling.

This paper studies whether a deliberately *small* temporal model — under 200K parameters, three orders of magnitude smaller than contemporary clinical foundation models — can satisfy all three properties simultaneously while remaining competitive with standard tabular baselines (logistic regression, gradient-boosted trees) on discrimination. We do not claim state-of-the-art AUROC; we claim that a small, calibrated, false-alarm-budgeted, externally-validated model is a more honest and more useful contribution than an incremental AUROC gain on a single internal split, and we provide the evidence to let a reader judge that trade-off directly.

1.1 Contributions

1. An open, fully reproducible pipeline (data ingestion, missingness-aware feature engineering, patient-level splitting, training, calibration, conformal risk control, ONNX export, and evaluation) for early

sepsis prediction on the PhysioNet Challenge 2019 dataset, requiring no institutional data use agreement and runnable end-to-end on an 8 GB-VRAM consumer GPU.

2. TinySepsis, a $<200\text{K}$ -parameter GRU architecture with an explicit missingness representation (value, observation mask, time-since-last-measurement, and first difference per channel), evaluated against zero-training clinical scores (qSOFA, NEWS2) and standard tabular ML baselines (logistic regression, XGBoost [Chen and Guestrin, 2016], LightGBM [Ke et al., 2017]).
3. A calibration and false-alarm-control pipeline combining temperature scaling, isotonic regression, and split-conformal risk control [Angelopoulos et al., 2022], evaluated for both calibration error and empirical false-alarm-rate coverage under real, uncurated distribution shift (hospital A $\rightarrow$ hospital B).
4. A cross-institution evaluation using the challenge’s own two-hospital split as a proxy — not a substitute — for the kind of independent-cohort external validation whose absence contributed to the well-documented failure of the Epic Sepsis Model [Wong et al., 2021], obtainable without credentialed MIMIC-IV/eICU access.
5. A full efficiency and deployability audit (parameter count, ONNX file size, CPU/GPU latency, a FastAPI research demo) so that "runs locally" is a measured claim, not an assertion.

1.2 What this paper is not

This is a methods and reproducibility contribution, not a clinical trial. TinySepsis has not been prospectively validated, is not FDA-cleared or CE-marked, and must not be used to make or influence real clinical decisions. All patient data originate from a fully de-identified, open-access research dataset; no new patient data were collected. We return to these limitations in Section 12.

2 Related Work

Sepsis prediction with machine learning. A substantial literature applies gradient boosting, recurrent networks, and Gaussian processes to early sepsis prediction on ICU time series, reviewed systematically by Moor et al. [2021] and, with a focus on diagnostic accuracy pooling, by Fleuren et al. [2020]. Representative systems include InSight/AISE [Nemati et al., 2018], a gradient-boosted model with a MIMIC-III external validation, and multi-output Gaussian process RNNs for real-time sepsis detection [Futoma et al., 2017]. The great majority of this literature reports AUROC/AUPRC as the primary (often sole) outcome and does not report calibration error, a false-alarm budget, or model size — the gap this paper targets.

The PhysioNet/CinC Challenge 2019. The dataset and task we use were introduced by Reyna et al. [2019b,a] as a public, unblinded competition with a novel time-dependent utility function rewarding early, penalizing late, true-positive predictions and penalizing false positives at a fixed rate. We adapt this utility function (Section 7) rather than reproducing the official multi-team leaderboard, since our study design deliberately restricts predictions to the pre-suspicion window (Section 3), which is not directly comparable to the original, full-stay scoring protocol.

Failures of external validation. The clearest cautionary tale for this literature is the Epic Sepsis Model: internally reported AUROC around 0.76–0.83 fell to 0.63 on independent, external evaluation, with a sensitivity of only 33% at the deployed threshold [Wong et al., 2021], prompting an editorial call for mandatory external validation of proprietary clinical AI [Habib et al., 2021]. We take this as the central methodological motivation for our hospital-A-train / hospital-B-test design.

Calibration and conformal risk control. Modern neural networks are frequently overconfident [Guo et al., 2017]; temperature scaling [Guo et al., 2017] and isotonic regression [Platt et al., 1999] are standard post-hoc fixes. Calibration is specifically flagged as underreported in clinical prediction modeling by Van Calster et al. [2019]. Conformal prediction [Vovk et al., 2005, Angelopoulos and Bates, 2021] provides finite-sample, distribution-free coverage guarantees; we use the conformal risk control framework of Angelopoulos et al. [2022] to select an alarm threshold with a controlled false-positive rate on exchangeable calibration data, and report how that guarantee degrades under genuine (non-exchangeable) hospital-to-hospital shift.

Missing data in clinical time series. Clinical measurement frequency is itself informative (informative missingness): a missing lactate is not the same event as a normal lactate. We follow the mask-and-time-since-last-measurement representation popularized by GRU-D and related work [Che et al., 2018], which is simpler than a full imputation-model but empirically competitive and requires no auxiliary generative model.

3 Problem Formulation

Cohort and time axis. Each ICU stay is represented as an irregularly-observed, hourly-binned multivariate time series over T hours (ICU length of stay, `ICULOS` in the source data), with vital signs, laboratory values, and static demographics recorded at each hour where available.

Clinical suspicion time. For a patient who develops sepsis, let t_{susp} denote the hour at which Sepsis-3 clinical criteria are met (organ dysfunction plus suspected infection, operationalized per Singer et al. [2016], Seymour et al. [2016]). The Challenge 2019 `SepsisLabel` column is constructed so that it equals 1 starting six hours before t_{susp} ; we recover t_{susp} per patient as $t_{\text{onset}} + 6$, where t_{onset} is the first hour with `SepsisLabel`=1.

Task. For a prediction horizon $W \in \{4, 6, 8\}$ hours, we define a binary label

$$y_W(t) = \mathbb{1} [0 < t_{\text{susp}} - t \leq W]$$

for septic patients, and $y_W(t) = 0$ for all t for patients who never meet Sepsis-3 criteria during their stay. Critically, we *truncate* each septic patient’s timeline at t_{susp} : no hour with $t \geq t_{\text{susp}}$ is used for training or evaluation. This is a deliberate, and we believe under-discussed, design choice: many published early-warning evaluations include post-suspicion hours as trivially-easy true positives, which inflates reported performance relative to the actual decision problem a clinician faces (predicting *before* recognition, not confirming after). Our primary horizon is $W = 6$ (matching the challenge’s native definition); $W = 4$ and $W = 8$ are secondary analyses reported as an ablation over horizon length.

Model input. At each eligible hour t , the model observes the trailing L -hour window $[\max(1, t - L + 1), t]$ (default $L = 24$, left-padded for early-stay hours) of vitals and labs, plus static age and sex, and must output a calibrated risk $\hat{p}_W(t) \in [0, 1]$ and, after thresholding, a binary alarm decision.

4 Dataset

We use the training partition of the PhysioNet/Computing in Cardiology Challenge 2019 dataset [Reyna et al., 2019b], comprising 40,336 ICU admissions across two hospital systems (Hospital A, $n = 20,336$;




Figure 1: TinySepsis architecture. Each hourly timestep is encoded as four channels per raw clinical feature (forward-filled value, observation mask, time-since-last-measurement, first difference), projected and passed through a small GRU whose initial hidden state is set from static demographics; the last non-padded timestep is read out through a compact MLP head. Total trainable parameters reported in Table 6.

Hospital B, $n = 20,000$), each with hourly vital signs (heart rate, oxygen saturation, temperature, systolic/mean/diastolic blood pressure, respiratory rate, end-tidal CO_2), 26 laboratory values (e.g. lactate, creatinine, white blood cell count, platelets, bilirubin), and static demographics (age, sex, ICU unit type, hours between hospital and ICU admission). The dataset is openly accessible without a PhysioNet credentialing process or data use agreement, which is precisely what makes it suitable for a fully reproducible, single-session study; we did not have credentialed access to MIMIC-IV [Johnson et al., 2023] or eICU [Pollard et al., 2018] at the time of this study, and we treat that as an explicit limitation (Section 11) rather than a gap to paper over.

Preprocessing. Physiologically implausible values (e.g. heart rate > 300 , pH outside $[6.5, 8.0]$) are set to missing. Patients with fewer than 4 recorded hours are dropped. All numeric features are forward-filled per patient; for each of the 34 time-varying features X we construct four channels: the forward-filled value, a binary observation mask (1 if X was actually measured this hour), the number of hours since the last real measurement (capped at 48h), and the first difference of the forward-filled value between consecutive hours. Leading (pre-first-observation) values are imputed with the training-set population mean at normalization time, not with zero, to avoid injecting a spurious, arbitrary reference point into a z-scored feature.

Splits. We split *by patient*, never by row, to prevent leakage of a single ICU stay across splits. Hospital A is stratified by ever-developing-sepsis and split 70/15/15 into train/validation/test. Hospital B is held out in its entirety as `external_test` and is never touched during training, model selection, or calibration fitting — it is used exactly once, for the numbers reported in Section 9.2. All normalization statistics (feature means/standard deviations) are fit exclusively on the Hospital A training split and frozen thereafter.

5 Model Architecture

TinySepsis (Figure 1) is a two-layer gated recurrent unit (GRU) network [Cho et al., 2014]. Each hourly timestep’s 136-dimensional input (34 features $\times$ 4 channels: value, mask, time-since-last-measurement, delta) is projected through a shared linear layer to a $H/2$ -dimensional token, then passed through a 2-layer GRU with hidden size $H = 128$. Static demographics (age, sex) are projected to initialize the GRU’s hidden state, rather than being concatenated at every timestep, which keeps the recurrent input dimension small. The GRU output at the last *real* (non-padded) timestep is passed through a LayerNorm $\rightarrow$ Linear $\rightarrow$ GELU $\rightarrow$ Dropout $\rightarrow$ Linear head producing a single risk logit. The resulting model has under 200K trainable parameters (exact count in Table 6), roughly two orders of magnitude smaller than the 1–5M-parameter target range typical of compact clinical time-series models, and four to five orders of magnitude smaller than general-purpose clinical language models.

Training. We use AdamW with cosine learning-rate annealing, automatic mixed precision (fp16 autocast) on GPU, gradient accumulation (2 steps, effective batch size 128), and class-weighted binary cross-entropy (positive weight set to the inverse empirical prevalence of the training split) to address the severe class imbalance intrinsic to hour-level sepsis prediction. Early stopping on validation AUROC (patience 4 epochs) selects the checkpoint used for all downstream calibration and evaluation.

6 Calibration and Conformal Risk Control

Raw sigmoid outputs from the trained GRU are calibrated in two stages, both fit exclusively on the validation split and then frozen:

1. **Temperature scaling.** A single scalar T is fit by minimizing validation-set BCE loss on logit/T [Guo et al., 2017], preserving the model’s ranking (AUROC is invariant to T) while correcting over/under-confidence.
2. **Isotonic regression.** A monotonic, non-parametric calibration map is fit on the temperature-scaled probabilities [Platt et al., 1999], correcting residual non-sigmoidal miscalibration that a single scalar cannot.

Conformal false-alarm-rate control. Independently of calibration, we select an alarm threshold τ using split-conformal risk control in the style of Angelopoulos et al. [2022]: on the validation split, τ is set to the $\lceil (n+1)(1-\alpha) \rceil / n$ quantile of the calibrated risk scores among *negative*-label hours, which guarantees (under the exchangeability of validation and future data) that the empirical false-alarm rate does not exceed α in expectation. We use $\alpha = 0.10$ as the primary operating point. Because Hospital B is, by construction, *not* exchangeable with Hospital A’s validation set (different institution, casemix, and equipment), we explicitly report how the empirical false-alarm rate on `external_test` deviates from the nominal α budget (Table 4) as a direct, quantitative measure of how much the conformal guarantee erodes under real-world distribution shift — a robustness check we consider at least as informative as the guarantee itself.

7 Clinical Utility Metric

In addition to standard discrimination and calibration metrics, we adapt the official Challenge 2019 time-dependent utility function [Reyna et al., 2019b], which assigns a reward or penalty to each hourly prediction as a piecewise-linear function of $\tau = t - t_{\text{susp}}$:

$$u_{\text{tp}}(\tau) = \begin{cases} \frac{1}{6}\tau + 2 & -12 \leq \tau \leq -6 \\ -\frac{1}{9}\tau + \frac{1}{3} & -6 < \tau \leq 3 \\ 0 & \text{otherwise} \end{cases} \quad u_{\text{fn}}(\tau) = \begin{cases} -\frac{2}{9}\tau - \frac{4}{3} & -6 < \tau \leq 3 \\ -2 & \tau > 3 \\ 0 & \tau \leq -6 \end{cases}$$

with $u_{\text{fp}} = -0.05$ and $u_{\text{tn}} = 0$ for every hour of a non-septic patient. We report the utility *normalized* against an oracle (perfect prediction) and a trivial “never alarm” policy, so that 1.0 corresponds to the oracle and 0.0 to inaction. We deviate from the official Challenge scoring protocol in one respect, stated explicitly to avoid any implied comparison to the public leaderboard: because TinySepsis only predicts on the pre-suspicion window (Section 3), the utility sum is computed only over that window, not over a patient’s full stay including post-suspicion hours. This is a narrower, and we believe more clinically honest, accounting of utility for a system designed as an early-warning tool rather than a continuous detector.

Table 1: Internal test set (Hospital A, held-out patients), primary 6h-horizon task.

Model	AUROC	AUPRC	Brier	ECE	Alarms/1000h	Norm. Utility
qSOFA (lite)	0.539	0.015	0.126	0.202	1000.0	-0.956
NEWS2 (lite)	0.590	0.018	0.063	0.165	810.6	-0.717
Logistic Regression	0.770	0.047	0.187	0.368	487.8	-0.013
XGBoost	0.807	0.058	0.117	0.257	417.1	0.133
LightGBM	0.782	0.058	0.113	0.237	475.6	0.007
TinySepsis (ours)	0.726	0.036	0.168	0.292	541.2	-0.141

8 Baselines

We compare TinySepsis against five baselines spanning the space from zero-training clinical heuristics to strong tabular gradient boosting:

- **qSOFA (lite)**. An approximation of quick-SOFA using the two of its three criteria available in this dataset (systolic BP ≤ 100 , respiratory rate ≥ 22); Glasgow Coma Scale is not recorded, so our qSOFA is a lower bound on the full 3-criterion score.
- **NEWS2 (lite)**. A National Early Warning Score 2 approximation [Royal College of Physicians, 2017] using the five vitals present in the dataset, omitting the supplemental-oxygen and consciousness sub-scores.
- **Logistic regression** on the same 136-dimensional current-hour feature vector as TinySepsis (no recurrence), class-balanced.
- **XGBoost** [Chen and Guestrin, 2016] and **LightGBM** [Ke et al., 2017], gradient-boosted trees on the identical tabular feature set, with `scale_pos_weight` set to the inverse training prevalence.

All learned baselines are trained on the same patient-level Hospital-A training split and evaluated on the identical test and external-test sets as TinySepsis, with no baseline-specific hyperparameter search beyond default-adjacent settings, to keep the comparison about architecture and representation rather than tuning budget.

9 Experiments and Results

9.1 Internal test set (Hospital A, held-out patients)

Table 1 reports discrimination (AUROC, AUPRC), calibration (Brier score, expected calibration error), operational alarm burden (alarms per 1000 patient-hours at the 85%-sensitivity operating point), and normalized clinical utility for all six models on the internal (same-hospital) test split.

Normalized utility is negative for four of the six models in Table 1, which is a direct consequence of the severe hour-level class imbalance (1.4% positive prevalence) rather than a broken metric. The alarms-per-1000h and utility columns are both computed at a shared, clinically-motivated 85%-sensitivity operating point, fixed across models for comparability; reaching 85% sensitivity on such a rare event necessarily requires a low decision threshold, so the majority of alarmed hours are false positives on non-septic patients. Each false alarm costs $u_{fp} = -0.05$ (Section 7), and because non-septic patient-hours vastly outnumber the modeled pre-suspicion septic-patient hours, this accumulated penalty exceeds the reward earned from correctly-flagged septic hours whenever a model’s false-alarm rate is high relative to its true-positive yield

Table 2: Cross-institution evaluation on Hospital B (different institution within the same data-collection effort, never used for training, validation, or calibration), 6h-horizon task.

Model	AUROC	AUPRC	Brier	ECE	Alarms/1000h	Norm. Utility
qSOFA (lite)	0.567	0.011	0.104	0.171	1000.0	-2.181
NEWS2 (lite)	0.642	0.014	0.059	0.159	782.2	-1.580
Logistic Regression	0.644	0.026	0.129	0.291	802.8	-1.678
XGBoost	0.687	0.029	0.075	0.178	716.3	-1.392
LightGBM	0.665	0.026	0.066	0.154	766.0	-1.558
TinySepsis (ours)	0.679	0.020	0.189	0.324	661.5	-1.205

— exactly the alarm-fatigue failure mode the utility metric is designed to surface [Sendak et al., 2020]. XGBoost is the exception (+0.133) precisely because it reaches the same sensitivity target with the lowest alarm rate in the table (417 vs. 604 alarms/1000h for TinySepsis), illustrating why alarms-per-1000h and normalized utility should always be read together rather than either in isolation.

9.2 Cross-institution evaluation (Hospital B)

Table 2 reports the identical metrics on Hospital B, a hospital system never seen during training, validation, calibration fitting, or model selection. We call this a cross-institution evaluation rather than an external validation throughout the rest of this paper: it is a genuine, uncurated distribution shift between two hospital systems, but both are part of the same PhysioNet Challenge 2019 collection effort, which is a weaker test than an independent-cohort external validation against a dataset assembled by a different team with different equipment and protocols (Limitations). The gap between Tables 1 and 2 is, by design, the primary robustness signal of this paper: a model whose AUROC or calibration collapses on Hospital B despite strong Hospital-A numbers reproduces exactly the failure mode documented for the Epic Sepsis Model [Wong et al., 2021], and a model that degrades gracefully is doing something more clinically meaningful than winning an internal leaderboard. Normalized utility is negative for *every* model here, including XGBoost: Hospital B’s lower sepsis prevalence (0.9% vs. 1.4%) shifts the same 85%-sensitivity threshold to an even less favorable false-positive-to-true-positive ratio, and the AUROC/calibration degradation documented above compounds the false-alarm-accumulation effect described for the internal split. This is the same mechanism, not a separate anomaly, and it is the clearest illustration in this paper of why a discrimination metric alone would have hidden a real deployability problem. It is also specific to the fixed 85%-sensitivity comparison threshold; Table 5 (Section 6) reports what happens at each model’s own conformal threshold instead, which changes this picture substantially and not entirely in TinySepsis’s favor.

9.3 Multi-seed replication

Tables 1–2 report a single training run per model, which is not sufficient evidence on its own for a claim about differential robustness (Section 11 discusses this directly). To address it, we retrained TinySepsis and all three learned baselines under 5 random seeds each (seeds 0–4; qSOFA/NEWS2 are deterministic scores and excluded), holding every other hyperparameter fixed, and recomputed the internal-test and external-test AUROC and their gap for each run. Table 3 reports the resulting mean $\pm$ standard deviation, and Figure 2 shows every individual run.

The gap is smaller for TinySepsis in every one of the 5 seeds than in any of the 15 tabular-baseline runs (5 seeds $\times$ 3 baselines): TinySepsis’s worst-case gap (seed 3, 0.089) is still below LightGBM’s best-case gap (seed 1, 0.105), and the two distributions do not overlap. A two-sided Welch’s t -test on the 5

Table 3: Internal AUROC, external AUROC, and their gap, mean $\pm$ std. over 5 random seeds. Logistic regression’s std is exactly 0 because scikit-learn’s `lbfgs` solver is deterministic given fixed data; its 5 runs are not 5 independent draws.

Model	Test AUROC	External AUROC	Gap	n
Logistic Regression	0.770 ± 0.000	0.644 ± 0.000	0.126 ± 0.000	5
XGBoost	0.805 ± 0.003	0.686 ± 0.003	0.119 ± 0.003	5
LightGBM	0.786 ± 0.002	0.670 ± 0.006	0.115 ± 0.007	5
TinySepsis (ours)	0.740 ± 0.010	0.667 ± 0.011	0.072 ± 0.015	5




Figure 2: Internal (left) and external (right) AUROC for all 5 seeds per model. Horizontal ticks mark the per-model mean. TinySepsis has visibly higher run-to-run variance than the tree ensembles, but its external AUROC does not collapse into the tabular baselines’ range in any of the 5 runs.

TinySepsis gaps against each baseline’s 5 gaps gives $p = 0.002$ (vs. XGBoost), $p = 0.002$ (vs. LightGBM), and $p = 0.001$ (vs. logistic regression) — all significant after a Bonferroni correction for the 3 comparisons ($\alpha = 0.05/3 = 0.0167$), with large effect sizes (Cohen’s $d \in [-4.9, -3.6]$). Because $n = 5$ per group is small, we corroborate this with the distribution-free one-sided Mann-Whitney U test, which gives $p = 0.004$ against all three baselines — the strongest signal $n = 5$ vs. $n = 5$ can produce, since it corresponds to complete rank separation between TinySepsis’s gaps and each baseline’s gaps. We report this as reasonably strong evidence, for this one hospital pair, that the headline finding is not an artifact of a favorably-chosen random seed. It remains evidence from a single hospital pair (Section 11); we make no claim about what a third institution, a different train/test split, or a different sepsis-onset definition would show.

9.4 Calibration and conformal risk control

Table 4 reports expected calibration error before and after temperature scaling and isotonic regression, and the empirical false-alarm rate achieved by the conformal threshold τ (fit on validation at target $\alpha = 0.10$) on each split, including the non-exchangeable external split.

Utility at the conformal threshold, all six models. Tables 1–2 deliberately use a fixed 85%-sensitivity operating point so all six models are compared under the same rule; that threshold is a comparison device, not the threshold this paper would recommend deploying. Table 5 instead reports alarm burden and

Table 4: Calibration error (ECE) before/after post-hoc calibration, and the conformal alarm threshold’s realized coverage.

Split	Raw ECE	Temp.-scaled ECE	Isotonic ECE	Conformal τ ($\alpha=0.10$)
Validation	0.297	0.311	0.000	0.026
Test (internal)	0.292	0.307	0.001	0.026
External (hosp. B)	0.324	0.335	0.005	0.026




Figure 3: Post-hoc-calibrated reliability (left) and clinical net benefit vs. treat-all/treat-none (right) for all six models on the internal test split.

normalized utility at each model’s *own* conformal threshold (fit per-model on validation, target $\alpha = 0.10$, same procedure as Section 6) — the operating point this paper actually argues for. The picture changes substantially: all four learned models (logistic regression, XGBoost, LightGBM, TinySepsis) achieve *positive* normalized utility internally, confirming that the negative numbers in Table 1 are an artifact of the fixed-sensitivity comparison threshold, not evidence that these models are unusable. Externally, however, a result we did not anticipate and report without softening: logistic regression, XGBoost, and LightGBM all remain utility-positive (0.12–0.16), while **TinySepsis is the only learned model that goes utility-negative on Hospital B** (−0.023, essentially break-even but on the wrong side of zero). The cause is visible in the realized-FPR column: TinySepsis’s conformal threshold, calibrated on Hospital-A validation data, controls the false-alarm rate to 10.1% internally (matching the 10% target almost exactly) but rises to 12.2% on Hospital B — roughly 2–3× the false-alarm rate of XGBoost (6.0%) or LightGBM (4.3%) on the same external split.¹ This is an important qualification to the paper’s headline finding (Section 10): TinySepsis’s *ranking* (AUROC) transfers across hospitals substantially better than the tabular baselines’, but its *calibrated deci-*

¹An earlier iteration of this analysis, run against a version of the model with a sequence-pooling implementation bug (since fixed; see `tests/test_model.py` for the regression test added and the code repository’s history for details), showed a substantially larger external false-alarm overshoot (~19.8% realized vs. the same 10% target). We report the corrected, current numbers throughout; we mention the discrepancy here in the interest of the same transparency this paper asks of the field.

Table 5: Alarm burden and normalized utility at each model’s own conformal threshold ($\alpha = 0.10$, fit on validation), all six models, internal test and external test.

Model	Split	τ	Realized FPR	Alarms/1000h	Norm. Utility
qSOFA (lite)	test	0.016	0.362	363.0	-0.264
qSOFA (lite)	external	0.016	0.311	311.8	-0.530
NEWS2 (lite)	test	0.022	0.060	60.9	-0.007
NEWS2 (lite)	external	0.022	0.061	62.1	-0.060
Logistic Regression	test	0.034	0.121	125.1	0.287
Logistic Regression	external	0.034	0.049	51.0	0.121
XGBoost	test	0.028	0.111	116.5	0.366
XGBoost	external	0.028	0.060	61.5	0.157
LightGBM	test	0.029	0.098	102.8	0.317
LightGBM	external	0.029	0.043	44.2	0.126
TinySepsis (ours)	test	0.026	0.101	103.8	0.137
TinySepsis (ours)	external	0.026	0.122	123.7	-0.023

Table 6: TinySepsis efficiency profile.

Metric	Value
Parameters	191,297
ONNX model size (MB)	0.77
PyTorch CPU latency (ms/sample)	1.976
ONNX Runtime CPU latency (ms/sample)	0.286
Peak training VRAM (GB)	0.07

sion threshold transfers worse than theirs, even though both effects are more modest than an earlier pass at this analysis suggested. A model can be the most robust choice on one axis and the least robust on another, and we did not know this until we computed both.

9.5 Efficiency and deployability

Table 6 reports the full efficiency profile used to substantiate the "runs locally" claim: parameter count, exported ONNX file size, and single-sample CPU/GPU inference latency, measured on the same consumer hardware (NVIDIA RTX 5060, 8GB VRAM) used for training.

9.6 Ablations

Table 7 isolates the contribution of (i) the explicit missingness representation (mask + time-since-last-measurement channels) and (ii) observation window length, by retraining TinySepsis with each component removed or altered and re-evaluating on both internal and external test sets.

Table 7: Ablation study: missingness representation and sequence length, TinySepsis (6h horizon).

Variant	Seq. len.	Test AUROC	External AUROC
Full (mask+tslm+delta)	24	0.726	0.680
No missingness encoding	24	0.689	0.694
Full, short sequence	12	0.717	0.661




Figure 4: Summed mean |SHAP value| per input-channel group (34 raw clinical variables each), XGBoost baseline, internal test split.

9.7 Feature attribution

To probe which of the four missingness-aware channels (measured value, observation mask, time-since-last-measurement, first difference) actually drive predictions rather than being included for symmetry, we compute TreeSHAP [Chen and Guestrin, 2016] attributions for the XGBoost baseline (which shares the identical 136-dimensional tabular feature set with TinySepsis) on a random 3000-row sample of the internal test split, and aggregate mean |SHAP| within each of the four channel groups across all 34 raw clinical variables. Figure 4 reports this channel-group importance and the top individual features; this is deliberately a coarser, more classical explainability analysis than the demo’s lightweight per-request attribution (Section 13), included here to support the missingness ablation in Table 7 with an independent, model-agnostic signal about which representation choices carry information.

9.8 Subgroup analysis

As a fairness and robustness probe (not a claim of fairness, following the framing of Obermeyer et al. [2019]), we stratify test-set AUROC by age band and recorded sex; results are reported in the accompanying results/tables/subgroup_analysis.csv and summarized in the repository README, since the sample size within some strata is too small for the point estimates to be reported in the main text with appropriate uncertainty.

10 Discussion

The central empirical question of this paper is not "does TinySepsis beat XGBoost on AUROC" — with under 200K parameters and no hyperparameter search, we did not expect it to dominate a well-tuned gradient-boosted tree on a tabular-friendly task, and on the internal test set it does not, in any of the 5 seeds we trained (Table 1, Table 3: XGBoost 0.805 ± 0.003 vs. TinySepsis 0.740 ± 0.010 mean AUROC). The result we consider load-bearing is the *gap*: every tabular baseline loses 0.115–0.126 AUROC on average moving from Hospital A to Hospital B, while TinySepsis loses only 0.072, and Section 9.3 shows this is not a lucky seed — TinySepsis’s degradation was smaller than all 15 baseline-seed runs, with $p \leq 0.002$ (Welch) and $p = 0.004$ (Mann-Whitney) against each baseline after Bonferroni correction. On the institution none of the models ever saw during development, TinySepsis trails XGBoost by a modest margin on average (mean external AUROC 0.667 vs. 0.686) and outperforms logistic regression (0.644) despite starting well behind both internally. We did not tune any model to produce this pattern — baselines used default-adjacent hyperparameters and the same feature set as TinySepsis (Section 8) — so the most parsimonious reading is that the higher-capacity tabular models are fitting institution-specific structure (Hospital A’s particular case mix, documentation habits, or equipment) that a smaller, more constrained recurrent model does not have the capacity to overfit to, at the cost of leaving internal AUROC on the table. This is exactly the failure mode reported for the Epic Sepsis Model in external deployment [Wong et al., 2021, Habib et al., 2021], reproduced here in miniature and under our control. We treat it as the paper’s main finding, not a secondary robustness footnote; it is now supported by more than a single run, but it remains evidence from a single hospital pair, and we do not know a priori whether it will replicate on a third institution (Section 11) — that replication is the most direct follow-up we can name. Two further comparisons matter beyond this headline result: (1) whether calibration and conformal risk control meaningfully change the alarms-per-1000-hours and normalized-utility numbers at a fixed sensitivity target, since those are the numbers that determine whether clinical staff tolerate a system long enough for it to matter [Sendak et al., 2020] — here isotonic regression collapses ECE by roughly two orders of magnitude (Table 4) while the conformal alarm-rate budget (target $\alpha = 0.10$) holds reasonably on the internal split but visibly inflates on Hospital B, a second, independent demonstration of the same cross-institution fragility; and (2) whether the missingness ablation (Table 7) shows that explicitly modeling *what wasn’t measured* carries real signal beyond the measured values themselves — removing the mask/time-since-last-measurement channels costs 0.037 AUROC internally, consistent with informative missingness being a first-class modeling concern in ICU time series rather than an implementation detail [Che et al., 2018], though its effect on external AUROC in our single-run comparison was noisier and should not be over-read (Section 11).

The headline finding needs one more qualification we did not anticipate when we designed the experiment: it holds for *ranking* (AUROC) but not for *calibrated decision thresholds*. Table 5 shows that TinySepsis’s own conformal threshold, fit on Hospital-A validation data exactly as prescribed in Section 6, controls the false-alarm rate almost exactly to the 10% target internally but rises to 12.2% on Hospital B — roughly 2–3 $\times$ the false-alarm rate XGBoost or LightGBM realize on the same external split with their own conformal thresholds. The net effect is that TinySepsis is the only learned model whose normalized utility goes *negative* on Hospital B at its recommended operating point (−0.023, essentially break-even), while the three tabular baselines all stay comfortably utility-positive there (+0.12 to +0.16). So the paper’s two central claims point in opposite directions depending on which axis of robustness is asked about: TinySepsis’s *relative ranking* of patients transfers across institutions better than the tabular baselines’, but its *absolute calibrated risk estimate* — the number a clinician would actually see, and the one the alarm threshold is set against — transfers worse. We do not have a confident explanation for this asymmetry; a plausible one is that a single global temperature/isotonic map is less able to correct a small recurrent model’s shift in output distribution under a change in base rate (0.9% vs. 1.4% prevalence) than it is for tree-ensemble scores, but we did not design an experiment to isolate this and flag it as the second highest-priority follow-up after the

multi-seed replication already discussed above. A reader taking one number away from this paper should not take "TinySepsis is the robust choice" – they should take "AUROC robustness and calibrated-threshold robustness are different properties that must both be measured, because a model can have one without the other."

11 Limitations

- **Single hospital pair, still.** All data originate from the PhysioNet Challenge 2019 release. Hospital A vs. Hospital B is a genuine, uncurated distribution shift, but it is one pair of institutions. Section 9.3 addresses *random-seed* variance directly: we retrained every learned model under 5 seeds and found TinySepsis’s internal-to-external AUROC gap smaller than all 15 baseline-seed runs, with $p \leq 0.002$ (Welch) and $p = 0.004$ (Mann-Whitney) against each baseline after Bonferroni correction, so the headline finding is not an artifact of a favorably-chosen initialization. What this does *not* address is *hospital-pair* variance: all 5 seeds share the same Hospital-A-train/Hospital-B-test split, so a systematic quirk of this specific pair of institutions (rather than of "internal vs. external" in general) would inflate every seed identically and our test would not catch it. We also note that logistic regression’s solver (`lbfgs`) is deterministic given fixed data, so its "5 seeds" are 5 copies of one run, not 5 independent draws – we still count it as $n = 5$ in the statistical tests above for consistency with the other baselines, which is conservative in TinySepsis’s favor only in the trivial sense that a zero-variance comparator cannot narrow the reported confidence; the effect sizes against XGBoost and LightGBM, which do have real seed-to-seed variance, are the ones that matter most. None of this is the same as validating against MIMIC-IV or eICU, which use different monitoring equipment, EHR systems, and documentation practices [Johnson et al., 2023, Pollard et al., 2018]; we did not have credentialed access to those datasets within the scope of this study, and we consider a MIMIC-IV/eICU replication – ideally with multiple hospital pairs, not just one more single pair – the single highest-priority follow-up.
- **Retrospective, not prospective.** All evaluation is retrospective on already-collected, already-labeled data. Retrospective AUROC and even retrospective alarm rates do not guarantee that a deployed system changes clinician behavior or patient outcomes; only a prospective, ideally randomized, evaluation can establish that, and none was performed here.
- **Approximate clinical scores.** Our qSOFA and NEWS2 baselines are necessarily incomplete (no Glasgow Coma Scale, no supplemental-oxygen flag in the source data), which likely understates their true achievable performance and should not be read as a definitive comparison against full-criteria clinical scores.
- **Label definition sensitivity.** We reconstruct t_{susp} from the Challenge’s own labeling convention rather than re-deriving Sepsis-3 criteria from raw data; sensitivity of our results to alternative sepsis-onset definitions is not evaluated here.
- **Conformal coverage under shift.** The conformal risk control guarantee is a marginal, exchangeability-dependent guarantee; Table 4 reports, rather than hides, its degradation on the external split, and that degradation should be treated as expected and diagnostic, not as a bug.
- **Demographic granularity.** The dataset provides only age and a binary sex field; no race, ethnicity, or insurance-status fields are available for a fuller equity audit, which limits the subgroup analysis in Section 9 to two axes.

12 Ethical Considerations

TinySepsis is a research artifact. It is **not** a validated medical device, has not undergone FDA clearance, CE marking, or any equivalent regulatory review, and must not be used to inform real clinical decisions. All underlying data are drawn from a fully de-identified, publicly released research dataset governed by PhysioNet’s open data license; no new patient data were collected, and no protected health information is processed by any code in this repository. The accompanying local demo (Section 13) displays an explicit "not for clinical use" disclaimer on every response and is intended solely to demonstrate that inference can run fully offline, on commodity hardware, without transmitting any data to a third party. Given the Epic Sepsis Model precedent [Wong et al., 2021, Habib et al., 2021], we consider it an ethical obligation — not merely good scientific practice — to report external-hospital degradation prominently rather than to lead with the more favorable internal numbers, and we have structured Section 9 accordingly. Any future clinical translation of this or a similar system would require, at minimum: prospective validation at the deployment site, a locally-fit conformal calibration (not a transplanted one), a human-in-the-loop workflow with an explicit override path, and institutional review board oversight.

13 Local Deployment and Reproducibility

The trained model is exported to ONNX (opset 17) and served through a minimal FastAPI application (`tinysepsis.demo.app`) that accepts hourly observations and returns a calibrated risk probability, a conformal alarm decision, and the top contributing (highest-magnitude, most-recently-observed) feature channels as a lightweight explanation, in the spirit of, but far simpler than, a full SHAP-based attribution. The demo runs entirely on CPU with no external network calls. The complete pipeline — dataset download, feature engineering, training, baseline training, calibration, conformal fitting, ONNX export, and evaluation — is implemented as a sequence of standalone scripts (`scripts/*.py`) intended to be run in order; exact commands, environment specification, and hardware requirements are documented in the repository `README.md`. All randomness-affecting steps use a fixed seed (Section 4) for patient-level splitting; full bitwise reproducibility of GPU-trained floating-point results is not claimed, consistent with known non-determinism in cuDNN/cuBLAS kernels.

14 Conclusion

We presented TinySepsis, a sub-200K-parameter recurrent model for early sepsis prediction designed around three properties we argue are under-weighted relative to raw discrimination in the sepsis-prediction literature: calibration, an explicit and auditable false-alarm budget via conformal risk control, and measured (not asserted) deployability on commodity hardware. Trained and evaluated entirely on the open-access PhysioNet Challenge 2019 dataset, with a cross-institution evaluation standing in – imperfectly, as discussed in Limitations – for the kind of independent-cohort evaluation whose absence contributed to a well-documented real-world failure [Wong et al., 2021], TinySepsis is offered as a fully reproducible baseline and a methodological argument: for clinical early-warning systems, how a model fails under distribution shift and how honestly its uncertainty is represented are not secondary concerns to be mentioned in a limitations paragraph, but primary evaluation criteria that belong in the main results table. Our own results make the case for that argument in a way we did not plan: TinySepsis’s ranking of patients (AUROC) generalized across hospitals substantially better than the tabular baselines’, but its calibrated alarm threshold generalized worse, and only measuring the first property would have produced a paper with the opposite conclusion from measuring both. We take that as the strongest evidence in this paper that discrimination alone is an insufficient basis for judging a clinical early-warning system – a conclusion that holds up under the 5-seed replication

in Section 9.3 ($p \leq 0.002$ against every tabular baseline), though it is still evidence from one hospital pair, and a multi-institution replication remains the necessary next step (Section 11).

Acknowledgments

This work uses the PhysioNet/Computing in Cardiology Challenge 2019 dataset [Reyna et al., 2019b], made openly available by its organizers and contributors; we thank them for enabling fully reproducible research on early sepsis prediction without a credentialing barrier.

Code and Data Availability

Code, trained model weights, the exact feature-engineering pipeline, and scripts to regenerate every table and figure in this paper from the raw, publicly downloadable dataset are available at <https://github.com/elianalfonsolopezpreciado/tinysepsis> (MIT license). The underlying clinical dataset is not redistributed; `scripts/download_data.py` fetches it directly from PhysioNet’s open-access S3 mirror.

References

- Anastasios N Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *arXiv preprint arXiv:2107.07511*, 2021.
- Anastasios N Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal risk control. *arXiv preprint arXiv:2208.02814*, 2022.
- Zhengping Che, Sanjay Purushotham, Kyunghyun Cho, David Sontag, and Yan Liu. Recurrent neural networks for multivariate time series with missing values. *Scientific Reports*, 8(1):6085, 2018. doi: 10.1038/s41598-018-24271-9.
- Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016. doi: 10.1145/2939672.2939785.
- Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using RNN encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- Lucas M Fleuren, Thomas LT Klausch, Charlotte L Zwager, Linda J Schoonmade, Tingjie Guo, Luca F Roggeveen, Eleonora L Swart, Armand RJ Girbes, Patrick Thorat, Ari Ercole, et al. Machine learning for the prediction of sepsis: A systematic review and meta-analysis of diagnostic test accuracy. *Intensive Care Medicine*, 46(3):383–400, 2020. doi: 10.1007/s00134-019-05872-y.
- Joseph Futoma, Sanjay Hariharan, and Katherine Heller. An improved multi-output Gaussian process RNN with real-time validation for early sepsis detection. *Proceedings of Machine Learning for Healthcare (MLHC)*, 2017.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, volume 70, pages 1321–1330, 2017.

- Andrew R Habib, Amy L Lin, and Ryan W Grant. The Epic sepsis model falls short—the importance of external validation. *JAMA Internal Medicine*, 181(8):1040–1041, 2021. doi: 10.1001/jamainternmed.2021.3333.
- Alistair EW Johnson, Lucas Bulgarelli, Lu Shen, Alvin Gayles, Ayad Shammout, Steven Horng, Tom J Pollard, Sicheng Hao, Benjamin Moody, Brian Gow, et al. MIMIC-IV, a freely accessible electronic health record dataset. *Scientific Data*, 10(1):1, 2023. doi: 10.1038/s41597-022-01899-x.
- Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Michael Moor, Bastian Rieck, Max Horn, Catherine R Jutzeler, and Karsten Borgwardt. Early prediction of sepsis in the ICU using machine learning: A systematic review. *Frontiers in Medicine*, 8:607952, 2021. doi: 10.3389/fmed.2021.607952.
- Shamim Nemati, Andre Holder, Fereshteh Razmi, Matthew D Stanley, Gari D Clifford, and Timothy G Buchman. An interpretable machine learning model for accurate prediction of sepsis in the ICU. *Critical Care Medicine*, 46(4):547–553, 2018. doi: 10.1097/CCM.0000000000002936.
- Ziad Obermeyer, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464):447–453, 2019. doi: 10.1126/science.aax2342.
- John Platt et al. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. *Advances in Large Margin Classifiers*, 10(3):61–74, 1999.
- Tom J Pollard, Alistair EW Johnson, Jesse D Raffa, Leo A Celi, Roger G Mark, and Omar Badawi. The eICU collaborative research database, a freely available multi-center database for critical care research. *Scientific Data*, 5:180178, 2018. doi: 10.1038/sdata.2018.178.
- Alvin Rajkomar, Jeffrey Dean, and Isaac Kohane. Machine learning in medicine. *New England Journal of Medicine*, 380(14):1347–1358, 2019. doi: 10.1056/NEJMr1814259.
- Matthew A Reyna, Chris Josef, Salman Seyedi, Russell Jeter, Supreeth P Shashikumar, M Brandon Westover, Ashish Sharma, Shamim Nemati, and Gari D Clifford. Early prediction of sepsis from clinical data: the PhysioNet/computing in cardiology challenge 2019. In *2019 Computing in Cardiology (CinC)*, pages 1–4. IEEE, 2019a. doi: 10.23919/CinC49843.2019.9005736.
- Matthew A Reyna, Christopher S Josef, Russell Jeter, Supreeth P Shashikumar, M Brandon Westover, Shamim Nemati, Gari D Clifford, and Ashish Sharma. Early prediction of sepsis from clinical data: The PhysioNet/computing in cardiology challenge. *Critical Care Medicine*, 48(2):210–217, 2019b. doi: 10.1097/CCM.0000000000004145.
- Chanu Rhee, Raymund Dantes, Lauren Epstein, David J Murphy, Christopher W Seymour, Theodore J Iwashyna, Sameer S Kadri, Derek C Angus, Robert L Danner, Anthony E Fiore, et al. Incidence and trends of sepsis in US hospitals using clinical vs claims data, 2009–2014. *JAMA*, 318(13):1241–1249, 2017. doi: 10.1001/jama.2017.13836.
- Chanu Rhee, Tracey M Jones, Yasin Hamad, Amila Pande, Jonathan Varon, Cameron O’Brien, Deverick J Anderson, David K Warren, Raymund B Dantes, Lauren Epstein, et al. Prevalence, underlying causes, and preventability of sepsis-associated mortality in US acute care hospitals. *JAMA Network Open*, 2(2):e187571, 2019. doi: 10.1001/jamanetworkopen.2018.7571.

Royal College of Physicians. National early warning score (NEWS) 2. *RCP London Report*, 2017.

Mark P Sendak, William Ratliff, Dina Sarro, Elizabeth Alderton, Joseph Futoma, Michael Gao, Marshall Nichols, Mike Revoir, Faraz Yashar, Christina Miller, et al. Real-world integration of a sepsis deep learning technology into routine clinical care: Implementation study. *JMIR Medical Informatics*, 8(7): e15182, 2020. doi: 10.2196/15182.

Christopher W Seymour, Vincent X Liu, Theodore J Iwashyna, Frank M Brunkhorst, Thomas D Rea, André Scherag, Gordon Rubenfeld, Jeremy M Kahn, Manu Shankar-Hari, Mervyn Singer, et al. Assessment of clinical criteria for sepsis: For the third international consensus definitions for sepsis and septic shock (Sepsis-3). *JAMA*, 315(8):762–774, 2016. doi: 10.1001/jama.2016.0288.

Christopher W Seymour, Foster Gesten, Hallie C Prescott, Marcus E Friedrich, Theodore J Iwashyna, Gary S Phillips, Stanley Lemeshow, Tiffany Osborn, Kathleen M Terry, and Mitchell M Levy. Time to treatment and mortality during mandated emergency care for sepsis. *New England Journal of Medicine*, 376(23): 2235–2244, 2017. doi: 10.1056/NEJMoa1703058.

Mervyn Singer, Clifford S Deutschman, Christopher W Seymour, Manu Shankar-Hari, Djillali Annane, Michael Bauer, Rinaldo Bellomo, Gordon R Bernard, Jean-Daniel Chiche, Craig M Coopersmith, et al. The third international consensus definitions for sepsis and septic shock (Sepsis-3). *JAMA*, 315(8):801–810, 2016. doi: 10.1001/jama.2016.0287.

Ben Van Calster, David J McLernon, Maarten Van Smeden, Laure Wynants, and Ewout W Steyerberg. Calibration: The achilles heel of predictive analytics. *BMC Medicine*, 17(1):230, 2019. doi: 10.1186/s12916-019-1466-7.

Vladimir Vovk, Alex Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.

Andrew Wong, Erkin Otles, John P Donnelly, Andrew Krumm, Jeffrey McCullough, Olivia DeTroyer-Cooley, Justin Pestrue, Marie Phillips, Judy Konye, Carleen Penzo, et al. External validation of a widely implemented proprietary sepsis prediction model in hospitalized patients. *JAMA Internal Medicine*, 181(8):1065–1070, 2021. doi: 10.1001/jamainternmed.2021.2626.